

Texture Quilting (Due Saturday 2/23/2019)

In this assignment, you will develop code to stitch together image patches sampled from an input texture in order to synthesize new texture images. You can download the test image used to generate the example above from assignment folder Canvas.

You should start by reading through the whole assignment, looking at the provided code in detail to make sure you understand what it does. The main function ***quilt_demo*** appears at the end. You will need to write several subroutines in order for it to function properly.

Name: Haochen Zhou

SID: 23567813

1. Shortest Path [25 pts]

Write a function ***shortest_path*** that takes an 2D array of ***costs***, of shape HxW, as input and finds the *shortest vertical path* from top to bottom through the array. A vertical path is specified by a single horizontal location for each row of the H rows. Locations in successive rows should not differ by more than 1 so that at each step the path either goes straight or moves at most one pixel to the left or right. The cost is the sum of the costs of each entry the path traverses. Your function should return an length H vector that contains the index of the path location (values in the range 0..W-1) for each of the H rows.

You should solve the problem by implementing the dynamic programming algorithm described in class. You will have a for-loop over the rows of the "cost-to-go" array (M in the slides), computing the cost of the shortest path up to that row using the recursive formula that depends on the costs-to-go for the previous row. Once you have get to the last row, you can then find the smallest total cost. To find the path which actually has this smallest cost, you will need to do backtracking. The easiest way to do this is to also store the index of whichever minimum was selected at each location. These indices will also be an HxW array. You can then backtrack through these indices, reading out the path.

Finally, you should create at least three test cases by hand where you know the shortest path and see that the code gives the correct answer.

In [1]:

```
1 #modules used in your code
2 import numpy as np
3 import matplotlib.pyplot as plt
```

In [2]:

```
1 def shortest_path(costs):
2     """
3     This function takes an array of costs and finds a shortest path from the
4     top to the bottom of the array which minimizes the total costs along the
5     path. The path should not shift more than 1 location left or right between
6     successive rows.
7
8     In other words, the returned path is one that minimizes the total cost:
9
10         total_cost = costs[0,path[0]] + costs[1,path[1]] + costs[2,path[2]] + .
11
12     subject to the constraint that:
13
14         abs(path[i]-path[i+1])<=1
15
16     Parameters
17     -----
18     costs : 2D float array of shape HxW
19             An array of cost values
```

```

20
21 Returns
22 -----
23 path : 1D array of length H
24         indices of a vertical path. path[i] contains the column index of
25         the path for each row i.
26 """
27
28 h,w=costs.shape
29 cost_arr = np.zeros((costs.shape))
30 index_arr = np.zeros((costs.shape))
31 path = np.zeros(h)
32
33 cost_arr[0]=costs[0]
34 index_arr[0]=np.arange(costs.shape[1])
35 for i in range(h):
36     if i<h-1:
37         #col edge
38         cost_arr[i+1][0]=costs[i+1][0] + np.min(cost_arr[i][:2])
39         index_arr[i+1][0]=cost_arr[i][:2].argmin()
40         cost_arr[i+1][-1]=costs[i+1][-1] + np.min(cost_arr[i][-2:])
41         index_arr[i+1][-1]= w-2+cost_arr[i][-2:].argmin()
42         #col mid
43         d = np.stack((cost_arr[i][:2],cost_arr[i][1:-1],cost_arr[i][2:]),a
44         cost_arr[i+1][1:-1]=costs[i+1][1:-1] + np.min(d,axis=1)
45         index_arr[i+1][1:-1]=np.argmin(d,axis=1)+np.arange(w-2)
46
47     #backtrack
48     index_arr = index_arr.astype(int)
49     path[-1]= np.argmin(cost_arr[-1])
50     m = np.argmin(cost_arr[-1])
51     for i in range(h-1,-1,-1):
52         if i>0:
53             path[i-1]=index_arr[i][m]
54             m = index_arr[i][m]
55
56     return path.astype(int)

```

In [3]:

```
1  #
2  # your test code goes here.  come up with at least 3 test cases
3  #
4  costs1 = np.array([[12,8,10,11],[10,15,7,17],[4,5,13,5]])
5  path1 = shortest_path(costs1)
6  print(costs1)
7  print(path1)
8
9  costs2 = np.array([[1,2,3,4,5],[6,7,8,9,10],[11,12,13,14,15]])
10 path2 = shortest_path(costs2)
11 print(costs2)
12 print(path2)
13
14 costs3 = np.array([[2,1,3,5],[2,11,2,12],[7,6,5,5],[2,1,3,5],[2,1,3,5],[2,1,3,5]])
15 path3 = shortest_path(costs3)
16 print(costs3)
17 print(path3)
```

```
[[12  8 10 11]
 [10 15  7 17]
 [ 4  5 13  5]]
[1 2 1]
[[ 1  2  3  4  5]
 [ 6  7  8  9 10]
 [11 12 13 14 15]]
[0 0 0]
[[ 2  1  3  5]
 [ 2 11  2 12]
 [ 7  6  5  5]
 [ 2  1  3  5]
 [ 2  1  3  5]
 [ 2  1  3  5]]
[1 2 2 1 1 1]
```

2. Image Stitching: [25 pts]

Write a function **stitch** that takes two gray-scale images, **left_image** and **right_image** and a specified **overlap** and returns a new output image by stitching them together along a seam where the two images have very similar brightness values. If the input images are of widths **w1** and **w2** then your stitched result image returned by the function should be of width **w1+w2-overlap** and have the same height as the two input images.



You will want to first extract the overlapping strips from the two input images and then compute a cost array given by the absolute value of their difference. You can then use your **shortest_path** function to find the seam along which to stitch the images where they differ the least in brightness. Finally you need to generate the output image by using pixels from the left image on the left side of the seam and from the right image on the right side of the seam. You may find it easiest to code this by first turning the path into an alpha mask for each image and then using the standard equation for compositing.

In [4]:

```
1  def stitch(left_image, right_image, overlap):
2      """
3      This function takes a pair of images with a specified overlap and stitches them
4      together by finding a minimal cost seam in the overlap region.
5
6      Parameters
7      -----
8      left_image : 2D float array of shape HxW1
9          Left image to stitch
10
11      right_image : 2D float array of shape HxW2
12          Right image to stitch
13
14      overlap : int
15          Width of the overlap zone between left and right image
16
17      Returns
18      -----
19      stitched : 2D float array of shape Hx(W1+W2-overlap)
20          The resulting stitched image
21      """
22      # inputs should be the same height
23      assert(left_image.shape[0]==right_image.shape[0])
24
25      h_left, w_left = left_image.shape
26      h_right, w_right = right_image.shape
27
28      path = shortest_path(abs(left_image[... ,w_left-overlap:] - right_image[... ,:w_left-overlap]))
29      stitched = np.zeros((h_left, (w_left+w_right-overlap)))
30      mask = np.zeros((h_left, overlap))
31      count = 0
32      for i in path:
33          mask[count][:i]=1
34          count+=1
35
36      stitched[... ,w_left-overlap:w_left]=mask*left_image[... ,w_left-overlap:] + (1-mask)*right_image[... ,:w_left-overlap]
37      stitched[... ,:w_left-overlap]=left_image[... ,:w_left-overlap]
38      stitched[... ,w_left:]=right_image[... ,overlap:]
39
40      assert(stitched.shape[0]==left_image.shape[0])
41      assert(stitched.shape[1]==(left_image.shape[1]+right_image.shape[1]-overlap))
42
43      return stitched
44
45
```

In [5]:

```
1 I = plt.imread("/Users/zhouhaochen/Desktop/assign3/rock_wall.jpg")
2 if (I.dtype == np.uint8):
3     I = I.astype(float) / 256
4
5 if (I.shape[-1]==3):
6     I = np.mean(I[:,:,:3],axis=-1)
7
8 plt.imshow(I,cmap=plt.cm.gray)
9 plt.show()
10
11 s = stitch(I,I,10)
12 s2 = stitch(I,s,10)
13 plt.imshow(s2,cmap=plt.cm.gray)
14 plt.show()
```



3. Texture Quilting: [25 pts]

Write a function ***synth_quilt*** that takes as input an array indicating the set of texture tiles to use, an array containing the set of available texture tiles, the ***tilesize*** and ***overlap*** parameters and synthesizes the output texture by stitching together the tiles. ***synth_quilt*** should utilize your stitch function repeatedly. First, for each horizontal row of tiles, construct the stitched row by repeatedly stitching the next tile in the row on to the right side of your row image. Once you have row images for all the rows, you can stitch them together to get the final image. Since your stitch function only works for vertical seams, you will want to transpose the rows, stitch them together, and then transpose the result. You may find it useful to look at the provided code below which simply puts down the tiles with the specified overlap but doesn't do stitching. Your quilting function will return a similar result but with much smoother transitions between the tiles.

In [6]:

```
1  def synth_quilt(tile_map,tiledb,tilesize,overlap):
2
3      """
4      This function takes as input an array indicating the set of texture tiles
5      to use at each location, an array containing the database of available texture
6      tiles, tilesize and overlap parameters, and synthesizes the output texture by
7      stitching together the tiles
8
9
10     Parameters
11     -----
12     tile_map : 2D array of int
13         Array storing the indices of which tiles to paste down at each output location
14
15     tiledb : 2D array of size ntiles x npixels
16         Collection of sample tiles to select from
17
18     tilesize : (int,int)
19         Size of a tile in pixels
20
21     overlap : int
22         Amount of overlap between tiles
23
24     Returns
25     -----
26     output : 2D float array
27         The resulting synthesized texture of size
28     """
29
30     # determine output size based on overlap and tile size
31     outh = (tilesize[0]-overlap)*tile_map.shape[0] + overlap
32     outw = (tilesize[1]-overlap)*tile_map.shape[1] + overlap
33     output = np.zeros((outh,outw))
34     col_vec0 = tiledb[tile_map[0,0],:]
35     col_vec0 = np.reshape(col_vec0,tilesize)
36
```

```

37     for i in range(tile_map.shape[0]):
38         tile_vec1 = tiledb[tile_map[i,0],:]
39         tile_image1 = np.reshape(tile_vec1,tilesizes)
40
41         for j in range(tile_map.shape[1]-1):
42             tile_vec2 = tiledb[tile_map[i,j+1],:]
43             tile_image2 = np.reshape(tile_vec2,tilesizes)
44             tile_image1 = stitch(tile_image1,tile_image2,overlap) #row
45
46             if i == 1:
47                 col0 = np.reshape(tiledb[tile_map[0,j+1],:],tilesizes)
48                 col_vec0 = stitch(col_vec0,col0,overlap)
49
50             if i>0:
51                 col_vec0 = stitch(np.transpose(col_vec0),np.transpose(tile_image1),c
52                 col_vec0 = np.transpose(col_vec0) #col
53
54
55     output[:,:] = col_vec0
56
57     return output
58

```

In [7]:

```

1  #I = plt.imread("/Users/zhouhaochen/Desktop/assign3/rock_wall.jpg")
2  #if (I.dtype == np.uint8):
3  #    I = I.astype(float) / 256
4  #if (I.shape[-1]==3):
5  #    I = np.mean(I[:,:,:3],axis=-1)
6
7  #tmap = np.zeros((2,2)).astype(int)
8  #im = np.array([I])
9  #sz = I.shape
10 #q = synth_quilt(tmap,im,sz,10)
11 #plt.imshow(sq,cmap=plt.cm.gray)
12 #plt.show()

```


4. Texture Synthesis Demo [25pts]

The function provided below **quilt_demo** puts together the pieces. It takes a sample texture image and a specified output size and uses the functions you've implemented previously to synthesize a new texture sample.

You should write some additional code in the cells that follow to in order demonstrate the final result and experiment with the algorithm parameters in order to produce a compelling visual result and write explanations of what you discovered.

Test your code on the provided image *rock_wall.jpg*. There are three parameters of the algorithm. The *tilesize*, *overlap* and *K*. In the provided `*texture_demo*` code below, these have been set at some default values. Include in your demo below images of three example texture outputs when you: (1) increase the tile size, (2) decrease the overlap, and (3) decrease the value for K. For each result explain how it differs from the default setting of the parameters and why.

Test your code on two other texture source images of your choice. You can use images from the web or take a picture of a texture yourself. You may need to resize or crop your input image to make sure that the **tiledb** is not overly large. You will also likely need to modify the **tilesize** and **overlap** parameters depending on your choice of texture. Once you have found good settings for these parameters, synthesize a nice output texture. Make sure you display both the image of the input sample and the output synthesis for your two other example textures in your submitted pdf.

In [8]:

```
1  #skimage is only needed for sample tiles code provided below
2  #you should not use it in your own code
3  import skimage as ski
4
5  def sample_tiles(image,tilesize,randomize=True):
6      """
7          This function generates a library of tiles of a specified size from a given
8
9          Parameters
10         -----
11         image : float array of shape HxW
12             Input image
13
14         tilesize : (int,int)
15             Dimensions of the tiles in pixels
16
17
18         Returns
19         -----
20         tiles : float array of shape  numtiles x numpixels
21             The library of tiles stored as vectors where npixels is the
22             product of the tile height and width
23         """
24
25  .....
```

```

25     tiles = ski.util.view_as_windows(image,tilesiz)
26     ntiles = tiles.shape[0]*tiles.shape[1]
27     npix = tiles.shape[2]*tiles.shape[3]
28     assert(npix==tilesiz[0]*tilesiz[1])
29
30     print("library has ntiles = ",ntiles,"each with npix = ",npix)
31
32     tiles = tiles.reshape((ntiles,npix))
33
34     # randomize tile order
35     if randomize:
36         tiles = tiles[np.random.permutation(ntiles),:]
37
38     return tiles
39
40
41 def topkmatch(tilestrip,dbstrips,k):
42     """
43     This function finds the top k candidate matches in dbstrips that
44     are most similar to the provided tile strip.
45
46     Parameters
47     -----
48     tilestrip : 1D float array of length npixels
49         Grayscale values of the query strip
50
51     dbstrips : 2D float array of size npixels x numtiles
52         Array containing brightness values of numtiles strips in the database
53         to match to the npixels brightness values in tilestrip
54
55     k : int
56         Number of top candidate matches to sample from
57
58     Returns
59     -----
60     matches : list of ints of length k
61         The indices of the k top matching tiles
62     """
63     assert(k>0)
64     assert(dbstrips.shape[0]>k)
65     error = (dbstrips-tilestrip)
66     ssd = np.sum(error*error,axis=1)
67     ind = np.argsort(ssd)
68     matches = ind[0:k]
69     return matches
70
71
72 def quilt_demo(sample_image, ntilesout=(10,20), tilesiz=(30,30), overlap=5, k=
73     """
74     This function takes an image and quilting parameters and synthesizes a
75     new texture image by stitching together sampled tiles from the source image
76
77
78     Parameters

```

```

79  -----
80  sample_image : 2D float array
81      Grayscale image containing sample texture
82
83  ntilesout : list of int
84      Dimensions of output in tiles, e.g. (3,4)
85
86  tileSize : int
87      Size of the square tile in pixels
88
89  overlap : int
90      Amount of overlap between tiles
91
92  k : int
93      Number of top candidate matches to sample from
94
95  Returns
96  -----
97  img : list of int of length K
98      The resulting synthesized texture of size
99  """
100
101  # generate database of tiles from sample
102  tiledb = sample_tiles(sample_image, tileSize)
103  # number of tiles in the database
104  nsampltiles = tiledb.shape[0]
105
106  if (nsampltiles < k):
107      print("Error: tile database is not big enough!")
108
109  # generate indices of the different tile strips
110  i, j = np.mgrid[0:tileSize[0], 0:tileSize[1]]
111  top_ind = np.ravel_multi_index(np.where(i < overlap), tileSize)
112  bottom_ind = np.ravel_multi_index(np.where(i >= tileSize[0] - overlap), tileSize)
113  left_ind = np.ravel_multi_index(np.where(j < overlap), tileSize)
114  right_ind = np.ravel_multi_index(np.where(j >= tileSize[1] - overlap), tileSize)
115
116  # initialize an array to store which tile will be placed
117  # in each location in the output image
118  tile_map = np.zeros(ntilesout, 'int')
119
120
121  #print('row:')
122  for i in range(ntilesout[0]):
123      #print(i)
124      for j in range(ntilesout[1]):
125
126          if (i==0) & (j==0): # first row first tile
127              matches = np.zeros(k) #range(nsampltiles)
128
129          elif (i==0): # first row (but not first tile)
130              left_tile = tile_map[i, j-1]
131              tilestrip = tiledb[left_tile, right_ind]
132              dbstrips = tiledb[:, left_ind]

```

```

133         matches = topkmatch(tilestrip,dbstrips,k)
134
135     elif (j==0):                                # first column (but not first row)
136         above_tile = tile_map[i-1,j]
137         tilestrip = tiledb[above_tile,bottom_ind]
138         dbstrips = tiledb[:,top_ind]
139         matches = topkmatch(tilestrip,dbstrips,k)
140
141     else:                                         # neighbors above and to the left
142         left_tile = tile_map[i,j-1]
143         tilestrip_1 = tiledb[left_tile,right_ind]
144         dbstrips_1 = tiledb[:,left_ind]
145         above_tile = tile_map[i-1,j]
146         tilestrip_2 = tiledb[above_tile,bottom_ind]
147         dbstrips_2 = tiledb[:,top_ind]
148         # concatenate the two strips
149         tilestrip = np.concatenate((tilestrip_1,tilestrip_2))
150         dbstrips = np.concatenate((dbstrips_1,dbstrips_2),axis=1)
151         matches = topkmatch(tilestrip,dbstrips,k)
152
153     #choose one of the k matches at random
154     tile_map[i,j] = matches[np.random.randint(0,k)]
155
156
157
158     output = synth_quilt(tile_map,tiledb,tilesize,overlap)
159
160     return output
161

```

In [9]:

```

1  # load in rock_wall.jpg
2  I = plt.imread("/Users/zhouhaochen/Desktop/assign3/rock_wall.jpg")
3  if (I.dtype == np.uint8):
4      I = I.astype(float) / 256
5  if (I.shape[-1]==3):
6      I = np.mean(I[:,:,:3],axis=-1)
7
8  # run and display results for quilt_demo with
9  #
10 # (0) default parameters
11 print("default")
12 np.random.seed(0)
13 q = quilt_demo(I)
14 plt.imshow(q,cmap=plt.cm.gray)
15 plt.show()
16
17 # (1) increased tile size
18 print("increase tile size")
19 np.random.seed(0)
20 q = quilt_demo(I, ntilesout=(10,20), tilesize=(100,100), overlap=5, k=5)
21 plt.imshow(q,cmap=plt.cm.gray)
22 plt.show()

```


```

23
24 # (2) decrease the overlap
25 print("decrease overlap")
26 np.random.seed(0)
27 q = quilt_demo(I, ntilesout=(10,20), tileSize=(30,30), overlap=2, k=5)
28 plt.imshow(q,cmap=plt.cm.gray)
29 plt.show()
30
31 # (3) increase the value for K.
32 print("increase k value")
33 np.random.seed(0)
34 q = quilt_demo(I, ntilesout=(10,20), tileSize=(30,30), overlap=5, k=100)
35 plt.imshow(q,cmap=plt.cm.gray)
36 plt.show()
37
38

```


default

library has ntiles = 29241 each with npix = 900



increase tile size


library has ntiles = 10201 each with npix = 10000




decrease overlap

library has ntiles = 29241 each with npix = 900





increase k value
 library has ntiles = 29241 each with npix = 900

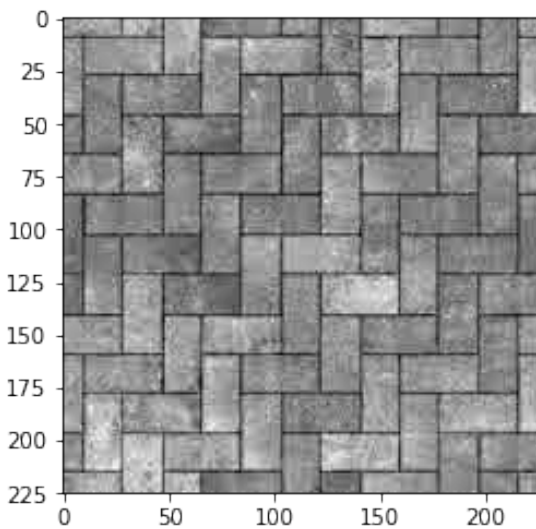


For each result shown, explain here how it differs visually from the default setting of the parameters and explain why:

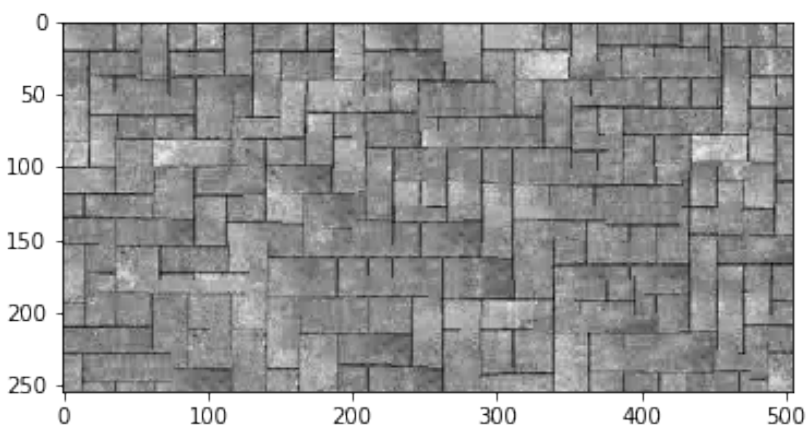
- . Increasing the tile size means that we will need more sampling images in the tile. And because the dimension of output in tiles will not change, the size of the sampling images will be smaller and the final texture image has to be shrunk. Thus, different with the default image, increasing the tile size will make the output image be denser.
- . Only decreasing the overlap will lead the dimension of the output image to be increased. And also there will be more sampling images in the output. This is because, decreasing the overlap means that the overlap region between left and right image will be decreased that results in a larger output image.
- . Increasing the value for k will make the result texture be more random and artificial. Increasing the k value means that there will be more tiles picked up from a list. Later, we have to choose one of those tiles at random. If the number of the tiles is large, later we will choose a tile more randomly that will lead to more errors and the texture to be artificial.

In [10]:

```
1  #
2  # load in yourimage1.jpg
3  #
4  # call quilt_demo, experiment with parameters if needed to get a good result
5  #
6  # display your source image and the resulting synthesized texture
7  #
8  I2 = plt.imread("/Users/zhouhaochen/Desktop/brick.jpg")
9  if (I2.dtype == np.uint8):
10     I2 = I2.astype(float) / 256
11  if (I2.shape[-1] == 3):
12     I2 = np.mean(I2[:,:,:3],axis=-1)
13
14  plt.imshow(I2,cmap=plt.cm.gray)
15  plt.show()
16
17  # run and display results for quilt_demo with
18  q = quilt_demo(I2)
19  plt.imshow(q,cmap=plt.cm.gray)
20  plt.show()
```




library has ntiles = 38416 each with npix = 900



In [11]:

```
1 #
2 # load in yourimage2.jpg
3 #
4 # call quilt_demo, experiment with parameters if needed to get a good result
5 #
6 # display your source image and the resulting synthesized texture
7 #
8 I3 = plt.imread("/Users/zhouhaochen/Desktop/stripes.jpg")
9 if (I3.dtype == np.uint8):
10     I3 = I3.astype(float) / 256
11 if (I3.shape[-1]==3):
12     I3 = np.mean(I3[:,:,:3],axis=-1)
13
14 plt.imshow(I3,cmap=plt.cm.gray)
15 plt.show()
16
17 # run and display results for quilt_demo with
18 q = quilt_demo(I3, ntilesout=(10,20), tilesiz=(60,60), overlap=10, k=5)
19 plt.imshow(q,cmap=plt.cm.gray)
20 plt.show()
```



library has ntiles = 5382 each with npix = 3600

