



Object Detection (due Saturday 3/9/2019)

In this assignment, you will develop an object detector based on gradient features and sliding window classification. A set of test images and *hogvis.py* are provided in the Canvas assignment directory

Name: Haochen Zhou

SID: 23567813

```
In [1]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
```

1. Image Gradients [20 pts]

Write a function that takes a grayscale image as input and returns two arrays the same size as the image, the first of which contains the magnitude of the image gradient at each pixel and the second containing the orientation.


Your function should filter the image with the simple x- and y-derivative filters described in class. Once you have the derivatives you can compute the orientation and magnitude of the gradient vector at each pixel. You should use ***scipy.ndimage.correlate*** with the 'nearest' option in order to nicely handle the image boundaries.


Include a visualization of the output of your gradient calculate for a small test image. For displaying the orientation result, please uses a cyclic colormap such as "hsv" or "twilight". (see <https://matplotlib.org/tutorials/colors/colormaps.html> (<https://matplotlib.org/tutorials/colors/colormaps.html>))

```
In [2]: 1  #we will only use:  scipy.ndimage.correlate
        2  from scipy import ndimage
        3
        4  def mygradient(image):
        5      """
        6          This function takes a grayscale image and returns two arrays of
        7          same size, one containing the magnitude of the gradient, the sec
        8          containing the orientation of the gradient.
        9
        10
        11      Parameters
        12      -----
        13      image : 2D float array of shape HxW
        14              An array containing pixel brightness values
        15
        16      Returns
        17      -----
        18      mag : 2D float array of shape HxW
        19              gradient magnitudes
        20
        21      ori : 2Dfloat array of shape HxW
        22              gradient orientations in radians
        23      """
        24      xk = np.array([[ -1, 1]])
        25      yk = np.array([[ -1], [ 1]])
        26
        27      # your code goes here
        28      mag = np.zeros(image.shape)
        29      ori = np.zeros(image.shape)
        30
        31      dx = ndimage.correlate(image,xk,mode="nearest")
        32      dy = ndimage.correlate(image,yk,mode="nearest")
        33
        34      mag = np.sqrt(np.square(dx) + np.square(dy))
        35      ori = np.arctan(dy/(dx+1.e-16))
        36
        37      return (mag,ori)
```

In [3]:

```
1 #
2 # Demonstrate your mygradient function here by loading in a grayscale
3 # image, calling mygradient, and visualizing the resulting magnitude
4 # and orientation images. For visualizing orientation image, I suggest
5 # using the hsv or twilight colormap.
6 #
7
8 [yy,xx] = np.mgrid[-40:40,-40:40]
9 image = np.array((xx*xx+yy*yy<=30*30),dtype=float)
10
11 (mag,ori) = mygradient(image)
12
13 #visualize results.
14 plt.imshow(mag,cmap=plt.cm.gray)
15 plt.show()
16
17 plt.imshow(ori,cmap=plt.cm.twilight)
18 plt.colorbar()
19 plt.show()
20
21
```





2. Histograms of Gradient Orientations [25 pts]

Write a function that computes gradient orientation histograms over each 8x8 block of pixels in an image. Your function should bin the orientation into 9 equal sized bins between $-\pi/2$ and $\pi/2$. The input of your function will be an image of size $H \times W$. The output should be a three-dimensional array **ohist** whose size is $(H/8) \times (W/8) \times 9$ where **ohist[i,j,k]** contains the count of how many edges of orientation k fell in block (i,j) . If the input image dimensions are not a multiple of 8, you should use **np.pad** with the **mode=edge** option to pad the width and height up to the nearest integer multiple of 8.

To determine if a pixel is an edge, we need to choose some threshold. I suggest using a threshold that is 10% of the maximum gradient magnitude in the image. Since each 8x8 block will contain a different number of edges, you should normalize the resulting histogram for each block to sum to 1 (i.e., **np.sum(ohist,axis=2)** should be 1 at every location).

I would suggest your function loops over the orientation bins. For each orientation bin you'll need to identify those pixels in the image whose magnitude is above the threshold and whose orientation falls in the given bin. You can do this easily in numpy using logical operations in order to generate an array the same size as the image that contains Trues at the locations of every edge pixel that falls in the given orientation bin and is above threshold. To collect up pixels in each 8x8 spatial block you can use the function **ski.util.view_as_windows(..., (8,8),step=8)** and **np.count_nonzeros** to count the number of edges in each block.

Test your code by creating a simple test image (e.g. a white disk on a black background), computing the descriptor and using the provided function **hogvis** to visualize it.

Note: in the discussion above I have assumed 8x8 block size and 9 orientations. In your code you should use the parameters **bsize** and **norient** in place of these constants.

```



In [4]: 1 #we will only use: ski.util.view_as_windows for computing hog descr
2 import skimage as ski
3 import math
4
5 def hog(image, bsize=8, norient=9):
6
7     """
8     This function takes a grayscale image and returns a 3D array
9     containing the histogram of gradient orientations descriptor (HOG)
10    We follow the convention that the histogram covers gradients sta
11    with the first bin at  $-\pi/2$  and the last bin ending at  $\pi/2$ .
12
13    Parameters
14    -----
15    image : 2D float array of shape HxW
16            An array containing pixel brightness values
17
18    bsize : int
19            The size of the spatial bins in pixels, defaults to 8
20
21    norient : int
22            The number of orientation histogram bins, defaults to 9
23
24    Returns
25    -----
26    ohist : 3D float array of shape (H/bsize,W/bsize,norient)
27            edge orientation histogram
28
29    """
30    # determine the size of the HOG descriptor
31    (h,w) = image.shape
32    h2 = int(np.ceil(h/float(bsize)))
33    w2 = int(np.ceil(w/float(bsize)))
34    ohist = np.zeros((h2,w2,norient))
35    # pad the input image on right and bottom as needed so that it
36    # is a multiple of bsize
37    pw = (0,int(bsize*w2-w))
38    ph = (0,int(bsize*h2-h))
39    image = np.pad(image, (ph,pw), 'symmetric')
40    # make sure we did the padding correctly
41    assert(image.shape==(h2*bsize,w2*bsize))
42
43    # compute image gradients
44    (mag,ori) = mygradient(image)
45    ori=-ori
46
47    # choose a threshold which is 10% of the maximum gradient magnit
48    thresh = 0.1*np.max(mag)
49
50    # separate out pixels into orientation channels, dividing the ra
51    #  $[-\pi/2,\pi/2]$  into norient equal sized bins and count how many
52    # as a sanity check, make sure every pixel gets assigned to at m
53    h = h2*bsize
54    w = w2*bsize
55
56    bin0 = -math.pi/2
57    bin_ = math.pi/norient
58    bincount = np.zeros((h2*bsize,w2*bsize))
59
60    for i in range(norient):
61        #create a binary image containing 1s for pixels at the ith

```

```
62         #orientation where the magnitude is above the threshold.
63         B = np.zeros((h,w))
64         B = (ori<=bin0+bin_)*(ori>=bin0)*(mag>thresh)
65         bin0=bin0+bin_
66
67         #sanity check
68         bincount = bincount + B
69         #pull out non-overlapping bsize x bsize blocks
70         chblock = ski.util.view_as_windows(B,(bsize,bsize),step=bsize)
71         #sum up the count for each block and store the results
72         sum_block = np.count_nonzero(chblock,axis=(-1,2))
73         ohist[:, :, i] = sum_block
74
75         assert(np.all(bincount<=1))
76         # lastly, normalize the histogram so that the sum along the orientation axis
77         # note: don't divide by 0! If there are no edges in a block (i.e
78         # is 0) then your code should leave all the values as zero.
79         k = np.sum(ohist,axis=2)
80         for i in range(norient):
81             ohist[:, :, i]=ohist[:, :, i]/(k+1.e-16)
82
83         assert(ohist.shape==(h2,w2,norient))
84         return ohist
```



```
In [5]: 1 #provided function for visualizing hog descriptors
2 import hogvis as hogvis
3
4 #
5 # generate a simple test image... a 80x80 image
6 # with a circle of radius 30 in the center
7 #
8 [yy,xx] = np.mgrid[-40:40,-40:40]
9 im = np.array((xx*xx+yy*yy<=30*30),dtype=float)
10 #
11 # display the image and the output of hogvis
12 #
13
14 plt.imshow(im)
15 plt.show()
16
17 h =hog(im,bsize=8,norient=9)
18 h2 = hogvis.hogvis(h,bsize=8,norient=9)
19
20 plt.imshow(h2)
21 plt.show()
22
```



3. Detection [25 pts]

Write a function that takes a template and an image and returns the top detections found in the image. Your function should follow the definition given below.

In your function you should first compute the histogram-of-gradient-orientation feature map for the image, then correlate the template with the feature map. Since the feature map and template are both three dimensional, you will want to filter each orientation separately and then sum up the results to get the final response. If the image of size $H \times W$ then this final response map will be of size $(H/8) \times (W/8)$.

When constructing the list of top detections, your code should implement non-maxima suppression so that it doesn't return overlapping detections. You can do this by sorting the responses in descending order of their score. Every time you add a detection to the list to return, check to make sure that the location of this detection is not too close to any of the detections already in the output list. You can estimate the overlap by computing the distance between a pair of detections and checking that the distance is greater than say 70% of the width of the template.

Your code should return the locations of the detections in terms of the original image pixel coordinates (so if your detector had a high response at block $[i,j]$ in the response map, then you should return $(8i,8j)$ as the pixel coordinates).

I have provided a function for visualizing the resulting detections which you can use to test your detect function. Please include some visualization of a simple test case.

```
In [6]: 1 #we will only use:  scipy.ndimage.correlate
2 from scipy import ndimage
3
4 def detect(image,template,ndetect=5,bsize=8,norient=9):
5
6     """
7     This function takes a grayscale image and a HOG template and
8     returns a list of detections where each detection consists
9     of a tuple containing the coordinates and score (x,y,score)
10
11     Parameters
12     -----
13     image : 2D float array of shape HxW
14             An array containing pixel brightness values
15
16     template : a 3D float array
17                 The HOG template we wish to match to the image
18
19     ndetect : int
20                 Number of detections to return
21
22     bsize : int
23                 The size of the spatial bins in pixels, defaults to 8
24
25     norient : int
26                 The number of orientation histogram bins, defaults to 9
27
28     Returns
```



```

29  -----
30  detections : a list of tuples of length ndetect
31      Each detection is a tuple (x,y,score)
32
33  """
34
35  # norient for the template should match the norient parameter pa
36  assert(template.shape[2]==norient)
37
38  fmap = hog(image,bsize=bsize,norient=norient)
39
40  #cross-correlate the template with the feature map to get the to
41  resp = np.zeros((fmap.shape[0],fmap.shape[1]))
42  for i in range(norient):
43      resp = resp + ndimage.correlate(fmap[:, :, i], template[:, :, i],
44
45  #sort the values in resp in descending order.
46  # val[i] should be ith largest score in resp
47  # ind[i] should be the index at which it occurred so that val[i]
48  val = np.sort(resp, axis=None)[::-1] #sorted response values
49  ind = np.argsort(resp.flatten())[::-1]
50
51  #work down the list of responses from high to low, to generate a
52  # list of ndetect top scoring matches which do not overlap
53  w=resp.shape[1]
54  detcount = 0
55  i = 0
56  detections = []
57  while ((detcount < ndetect) and (i < len(val))):
58      # convert 1d index into 2d index
59      yb = ind[i]//w
60      xb = ind[i]%w
61      assert(val[i]==resp[yb,xb]) #make sure we did indexing corre
62
63      #covert block index to pixel coordinates based on bsize
64      xp = xb*bsize
65      yp = yb*bsize
66
67      #check if this detection overlaps any detections that we've
68      #to the list. compare the x,y coordinates of this detection
69      #coordinates of the detections already in the list and see i
70      #by checking if the distance between them is less than 70% o
71      # width/height
72      y = [detections[y][1] for y in range(len(detections))]
73      x = [detections[x][0] for x in range(len(detections))]
74      overlap = np.any(np.sqrt(np.square(xp-x)+np.square(yp-y)) <
75
76      #if the detection doesn't overlap then add it to the list
77      if not overlap:
78          detcount = detcount + 1
79          detections.append((xp,yp,val[i]))
80      i=i+1
81
82  if (len(detections) < ndetect):
83      print('WARNING: unable to find ',ndetect,' non-overlapping d
84
85  return detections

```

```

In [7]: 1 import matplotlib.patches as patches
2
3 def plot_detections(image,detections,tsize_pix):
4     """
5     This is a utility function for visualization that takes an image
6     a list of detections and plots the detections overlayed on the image
7     as boxes.
8
9     Color of the bounding box is based on the order of the detection
10    the list, fading from green to red.
11
12    Parameters
13    -----
14    image : 2D float array of shape HxW
15            An array containing pixel brightness values
16
17    detections : a list of tuples of length ndetect
18                 Detections are tuples (x,y,score)
19
20    tsize_pix : (int,int)
21                The height and width of the box in pixels
22
23    Returns
24    -----
25    None
26
27    """
28    ndetections = len(detections)
29
30    plt.imshow(image)
31    ax = plt.gca()
32    w = tsize_pix[1]
33    h = tsize_pix[0]
34    red = np.array([1,0,0])
35    green = np.array([0,1,0])
36    ct = 0
37    for (x,y,score) in detections:
38        xc = x-(w//2)
39        yc = y-(h//2)
40        col = (ct/ndetections)*red + (1-(ct/ndetections))*green
41        rect = patches.Rectangle((xc,yc),w,h,linewidth=3,edgecolor=col)
42        ax.add_patch(rect)
43        ct = ct + 1
44
45    plt.show()


```

```

In [8]: 1 #
2 # sketch of some simple test code, modify as needed
3 #
4
5 #create a synthetic image
6 [yy,xx] = np.mgrid[-40:40,-40:40]
7 im1 = np.array((xx*xx+yy*yy<=30*30),dtype=float)
8 [yy,xx] = np.mgrid[-60:20,-60:20]
9 im2 = np.array((xx*xx+yy*yy<=25*25),dtype=float)
10 im = 0.5*im1+0.5*im2
11
12
13 #compute feature map with default parameters
14 fmap = hof(im)

```

```
15
16 #extract a 3x3 template
17 template = fmap[1:4,1:4,:]
18
19 plt.imshow(im[8:32,8:32])
20 plt.show()
21
22 #run the detect code
23 detections = detect(im,template,ndetect=5)
24
25 #visualize results.
26 plot_detections(im,detections,(24,24))
27
28 # visually confirm that:
29 # 1. top detection should be the same as the location where we sel
30 # 2. multiple detections do not overlap too much
```



4. Learning Templates [15 pts]

The final step is to implement a function to learn a template from positive and negative examples. Your code should take a collection of cropped positive and negative examples of the object you are interested in detecting, extract the features for each, and generate a template by taking the average positive template minus the average negative template.

```
In [9]: 1 def learn_template(posfiles,negfiles,tsize=np.array([16,16]),bsize=8
```

```

2      """
3      This function takes a list of positive images that contain croppe
4      examples of an object + negative files containing cropped backgr
5      and a template size. It produces a HOG template and generates vi
6      of the examples and template
7
8      Parameters
9      -----
10     posfiles : list of str
11                Image files containing cropped positive examples
12
13     negfiles : list of str
14                Image files containing cropped negative examples
15
16     tsize : (int,int)
17            The height and width of the template in blocks
18
19     Returns
20     -----
21     template : float array of size tsize x norient
22                The learned HOG template
23
24     """
25
26     #compute the template size in pixels
27     #corresponding to the specified template size (given in blocks)
28     tsize_pix=bsize*tsize
29
30     #figure to show positive training examples
31     pltct = 1
32
33     #accumulate average positive and negative templates
34     pos_t = np.zeros((tsize[0],tsize[1],norient),dtype=float)
35     for file in posfiles:
36         #load in a cropped positive example
37         img = plt.imread(file)
38         if (img.dtype == np.uint8):
39             img = img.astype(float) / 256
40
41         #convert to grayscale and resize to fixed dimension tsize_pi
42         #using skimage.transform.resize if needed.
43         if (img.shape[-1]==3):
44             img = np.mean(img[:,:,:3],axis=-1)
45
46         img = ski.transform.resize(img,(tsize[0]*bsize,tsize[1]*bsiz
47
48         #display the example. if you want to train with a large # of
49         #you may want to modify this, e.g. to show only the first 5.
50         plt.imshow(img,cmap=plt.cm.gray)
51         plt.show()
52
53         #extract feature
54         fmap = hog(img,bsize,norient)
55
56         f = hogvis.hogvis(fmap,bsize=8,norient=9)
57         plt.imshow(f)
58         plt.show()
59
60         #compute running average
61         pos_t = pos_t + fmap[:tsize[0],:tsize[1],:]
62

```

```
63 pos_t = (1/len(posfiles))*pos_t
64
65 # repeat same process for negative examples
66 neg_t = np.zeros((tsize[0],tsize[1],norient),dtype=float)
67 for file in negfiles:
68
69     img = plt.imread(file)
70     if (img.dtype == np.uint8):
71         img = img.astype(float) / 256
72
73     if (img.shape[-1]==3):
74         img = np.mean(img[:,:,:3],axis=-1)
75     img = ski.transform.resize(img,(tsize[0]*bsize,tsize[1]*bsiz
76
77     plt.imshow(img,cmap=plt.cm.gray)
78     plt.show()
79
80     fmap = hog(img,bsize,norient)
81     f = hogvis.hogvis(fmap,bsize=8,norient=9)
82     plt.imshow(f)
83     plt.show()
84
85     neg_t = neg_t+fmap[:tsize[0],:tsize[1],:]
86
87 neg_t = (1/len(negfiles))*neg_t
88
89 # now construct our template as the average positive minus avera
90 template = pos_t - neg_t
91
92
93 return template
94
```

In []:

1

5. Experiments [15 pts]

Test your detection by training a template and running it on a test image.

In your experiments and writeup below you should include: (a) a visualization of the positive and negative patches you use to train the template and corresponding hog feature, (b) the detection results on the test image. You should show (a) and (b) for **two different object categories**, the provided face test images and another category of your choosing (e.g. feel free to experiment with detecting cat faces, hands, cups, chairs or some other type of object). Additionally, please include results of testing your detector where there are at least 3 objects to detect (this could be either 3 test images which each have one or more objects, or a single image with many (more than 3) objects). Your test image(s) should be distinct from your training examples. Finally, write a brief (1 paragraph) discussion of where the detector works well and when it fails. Describe some ways you might be able to make it better.

NOTE 1: You will need to create the cropped test examples to pass to your **learn_template**. You can do this by cropping out the examples by hand (e.g. using an image editing tool). You should attempt to crop them out in the most consistent way possible, making sure that each example is centered with the same size and aspect ratio. Negative examples can be image patches that don't contain the object of interest. You should crop out negative examples with roughly the same resolution as the positive examples.

NOTE 2: For the best result, you will want to test on images where the object is the same size as your template. I recommend using the default **bsize** and **norient** parameters for all your experiments. You will likely want to modify the template size as needed

Experiment 1: Face detection



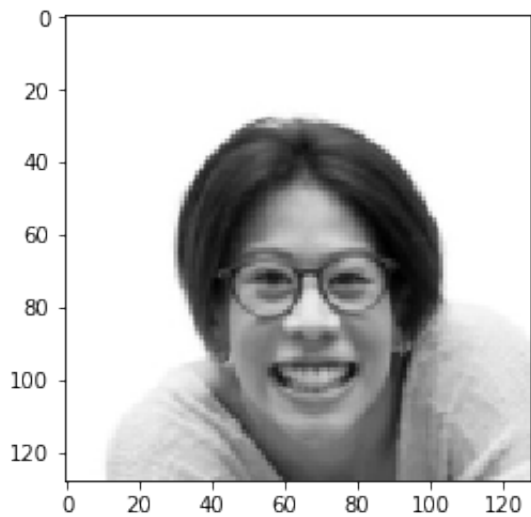
```
In [10]: 1 # assume template is 16x16 blocks, you may want to adjust this
2 # for objects of different size or aspect ratio.
3 # compute image a template size
4 bsize=8
5 tsize=np.array([16,16]) #height and width in blocks
6 tsize_pix = bsize*tsize #height and width in pixels
7
8 path="/Users/zhouhaochen/Desktop/assign4/assignment4_files/images/fac
9 path2="/Users/zhouhaochen/Desktop/assign4/assignment4_files/images/fa
10
11 posfiles = (path+'f1.jpg',path+'f2.jpg',path+"f3.jpg",path+"f4.jpg",
12 negfiles = (path2+'b1.jpg',path2+'b2.jpg',path2+'b3.jpg')
13
14 # call learn_template to learn and visualize the template and traini
15 template = learn_template(posfiles,negfiles,tsize=tsize)
16
17 # call detect on one or more test images, visualizing the result wit
18 I = plt.imread("/Users/zhouhaochen/Desktop/assign4/assignment4_files/
19
20 if (I.dtype == np.uint8):
21     I = I.astype(float) / 256
22 if (I.shape[-1]==3):
23     I = np.mean(I[:,:,:3],axis=-1)
24
```

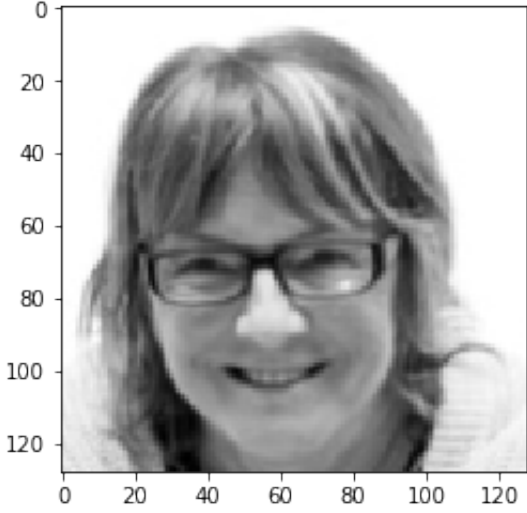


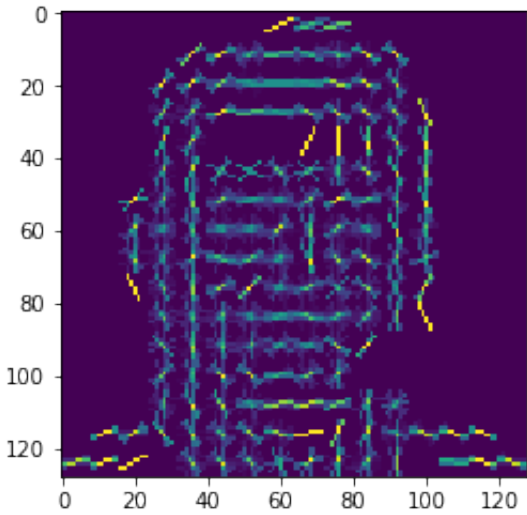






```
25 detections = detect(I,template,ndetect=4,bsize=8,norient=9)
26 plot_detections(I,detections,tsize_pix)
27
```





/Users/zhouhaochen/anaconda3/lib/python3.7/site-packages/skimage/transform/_warps.py:105: UserWarning: The default mode, 'constant', will be changed to 'reflect' in skimage 0.15.
warn("The default mode, 'constant', will be changed to 'reflect' in "

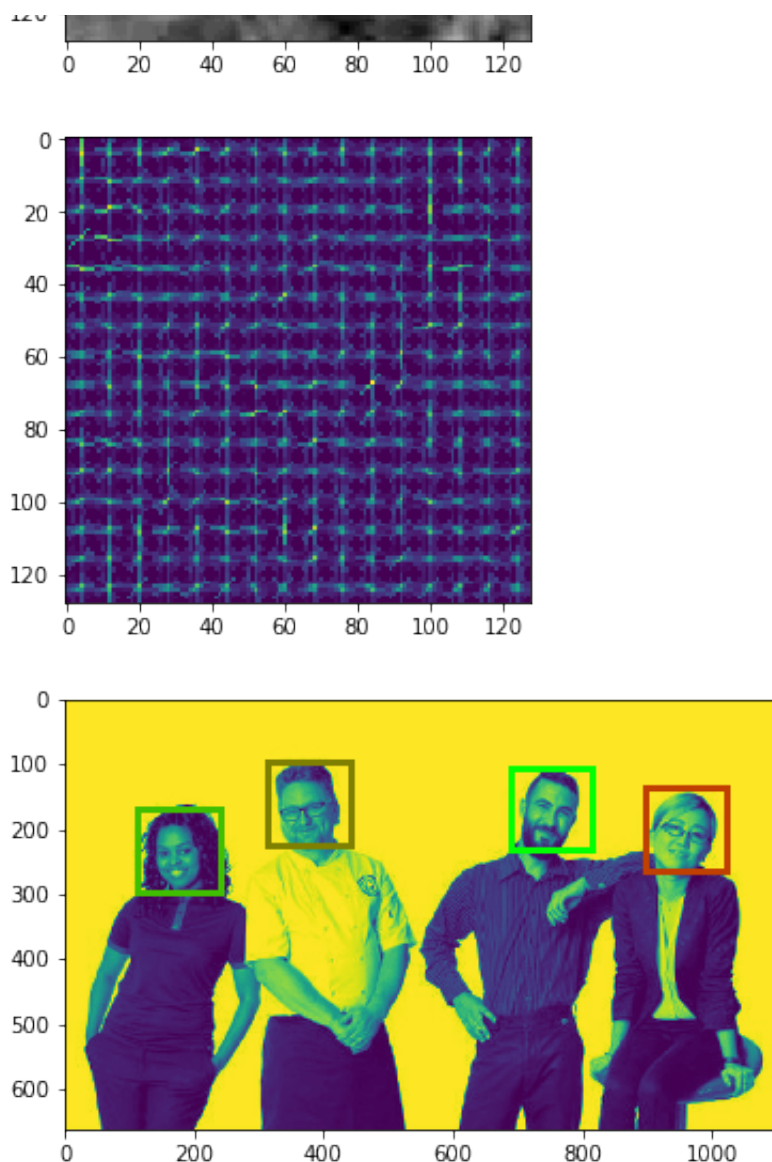
/Users/zhouhaochen/anaconda3/lib/python3.7/site-packages/skimage/transform/_warps.py:110: UserWarning: Anti-aliasing will be enabled by default in skimage 0.15 to avoid aliasing artifacts when down-sampling images.
warn("Anti-aliasing will be enabled by default in skimage 0.15 to "








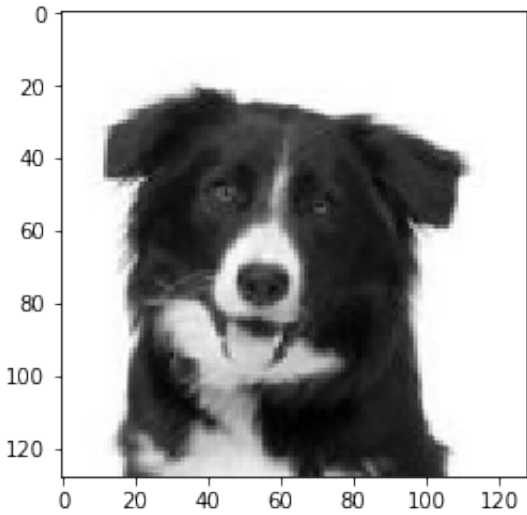



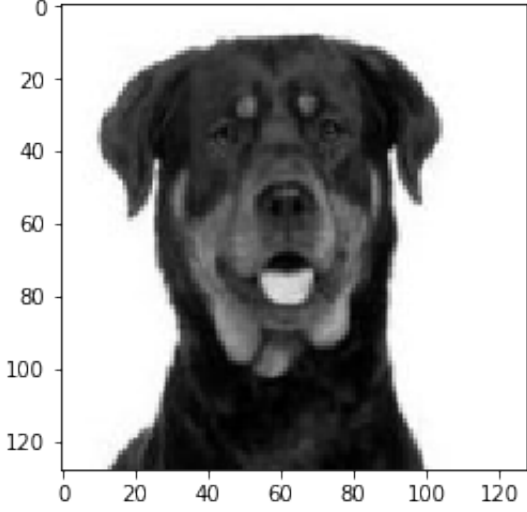

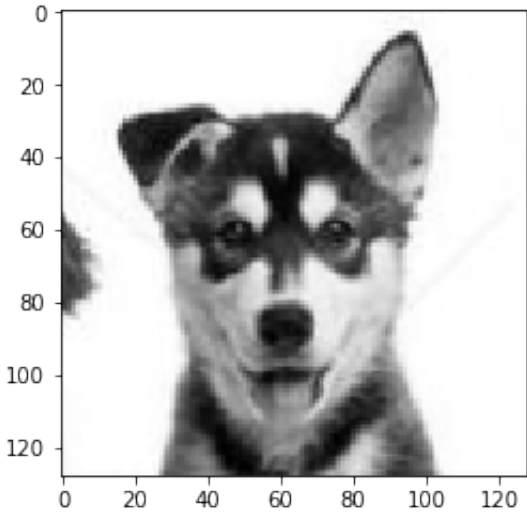









Experiment 2: ??? detection

```
In [11]: 1 # assume template is 16x16 blocks, you may want to adjust this
2 # for objects of different size or aspect ratio.
3 # compute image a template size
4 bsize=8
5 tsize=np.array([16,16]) #height and width in blocks
6 tsize_pix = bsize*tsize #height and width in pixels
7 path="/Users/zhouhaochen/Desktop/assign4/assignment4_files/images/dog
8 path2="/Users/zhouhaochen/Desktop/assign4/assignment4_files/images/do
9 posfiles = (path+"d1.jpg",path+"d2.jpg",path+"d3.jpg",path+"d4.jpg",
10 negfiles = (path2+'b1.jpg',path2+'b1.jpg')
11
12 # call learn_template to learn and visualize the template and traini
13 template = learn_template(posfiles,negfiles,tsize=tsize)
14
15 # call detect on one or more test images, visualizing the result wit
16 I = plt.imread("/Users/zhouhaochen/Desktop/assign4/assignment4_files/
17 if (I.dtype == np.uint8):
18     I = I.astype(float) / 256
19 if (I.shape[-1]==3):
20     I = np.mean(I[:,:,:3],axis=-1)
21
22 detections = detect(I,template,ndetect=6,bsize=8,norient=9)
23 plot_detections(I,detections,tsize_pix)
```









```
In [ ]: 1
```

If the size of the detected objects in the test image are too small or large when comparing to the size of the bounding box, the detector will fail. For example, some dogs' faces are really small in the test image and cannot be caught by the bounding box correctly. If the size of the bounding box can be automatically changed with the size of the detected objects, the result will be better. In addition, if the positive image has a different orientation with the detected image, the detector will also fail. If they have the same orientation, that will lead to a good result. We may need more positive images or templates to train our detector.

```
In [ ]: 1
```

